

Jayant Sharma

+91-9390126173 | mejayantsharma130@gmail.com | [LinkedIn](#) | [GitHub](#) | Hyderabad, India

SUMMARY

Final-year B.Tech student in Computer Science and Data Science, building applied ML and LLM-agent projects end to end — from data pipeline through statistical modeling to a deployed app. Comfortable across classical ML (scikit-learn, XGBoost), explainability (SHAP), and LLM tooling. Looking for a Data Science, ML, or AI internship.

PROJECTS

AI Risk Manager — Razorpay Internship Buildathon (Track 2) [GitHub](#)

XGBoost, SHAP, entity risk memory, LLM explainer agent · built in a 5-day buildathon

- Built an entity-aware transaction risk system on the IEEE-CIS fraud dataset: an XGBoost model scores each transaction, SHAP explains the top contributing factors, and an LLM agent turns that into a plain-English explanation — with the LLM kept off the decision path entirely.
- Designed a rolling, severity-weighted risk memory per entity (card/account) that escalates repeat offenders into WATCH/ELEVATED status, with cutoffs chosen by a cost-based grid sweep rather than a fixed rule.
- Picked decision thresholds by minimizing an explicit cost function (false negatives × avg. fraud loss + false positives × avg. false-positive cost) instead of optimizing for AUC alone, with a sensitivity sweep over plausible cost assumptions.

lazy-rag — Adaptive Retrieval for RAG [GitHub](#)

Independent research project · benchmarked against Self-RAG / FLARE / Adaptive-RAG baselines

- Tested whether a model's self-reported confidence can trigger retrieval only when it's actually needed, instead of retrieving on every query, benchmarking no-RAG, always-RAG, and a confidence-gated "lazy-RAG" pipeline on PopQA.
- Found the negative result and reported it honestly: answer quality stayed roughly flat while retrieval rate crept up as the confidence threshold loosened — self-verbalized confidence is a weak retrieval trigger, evidence against a naive version of the approach.

Life Expectancy Predictor [GitHub](#)

scikit-learn, Streamlit

- Trained a Random Forest Regressor on WHO life-expectancy data (adult mortality, GDP, schooling, immunization) reaching $R^2 = 0.968$ and MAE = 1.08 years on held-out data, and shipped it as an interactive Streamlit app.

LEADERSHIP & EXTRACURRICULARS

Secretary — IEEE Systems, Man & Cybernetics Society

Jan – Apr 2025

GNITC, Hyderabad

- Planned and ran technical workshops end to end — scheduling, communications, and day-of logistics.
- Maintained meeting minutes, official correspondence, and chapter documentation.

EDUCATION

B.Tech, Computer Science (Data Science) — Gurunanak Institutions Technical Campus, Hyderabad 2023 – 2027 (Expected)

Higher Secondary Certificate — Sri Chaitanya Jr. Kalasala, Hyderabad

2023

SKILLS & CERTIFICATIONS

Technical: Python, Pandas, NumPy, scikit-learn, XGBoost, SHAP, Statistical Analysis, FastAPI, Streamlit, SQL, LLM/Prompt Engineering, Git/GitHub

Certifications: IBM Data Science Professional Certificate (Coursera) · Data Science for Engineers (NPTEL / IIT)

Languages: English (C1), Hindi (C2)